

Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>



deeplearning.ai

Introduction to ML strategy

Why ML
Strategy?

Motivating example



90%

Ideas:

- Collect more data ←
- Collect more diverse training set
- Train algorithm longer with gradient descent
- Try Adam instead of gradient descent
- Try bigger network
- Try smaller network
- Try dropout
- Add L_2 regularization
- Network architecture
 - Activation functions
 - # hidden units
 - ...




deeplearning.ai


Introduction to ML strategy

Orthogonalization

TV tuning example



Orthogonalization



Car



→ Steering]

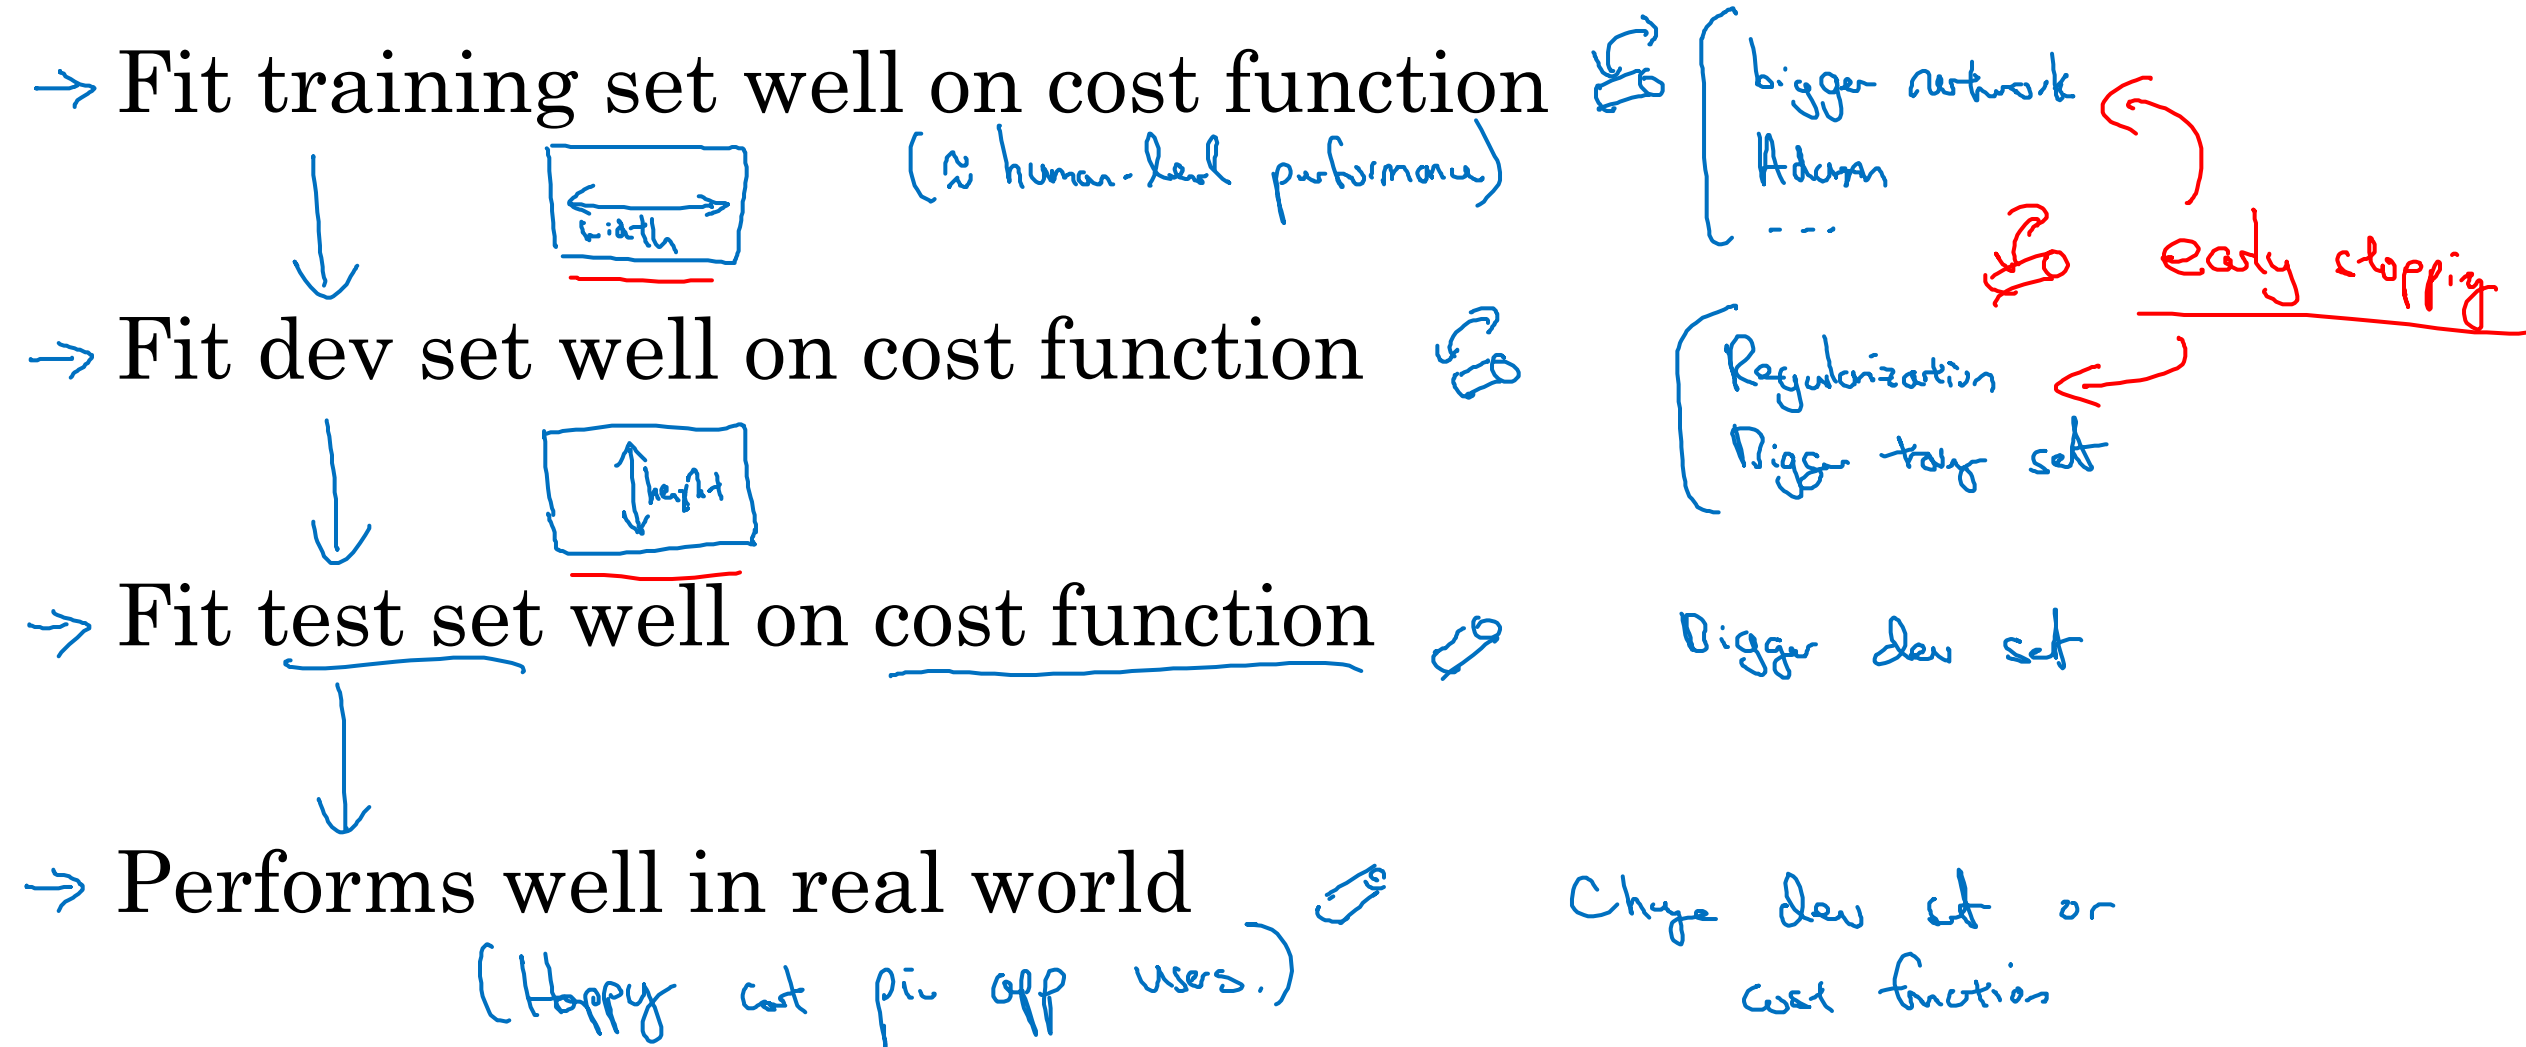
→ { Accelerator
Braking }

→ $0.3 \times \text{angle} - 0.8 \text{ speed}$

→ $2 \times \text{angle} + 0.9 \text{ speed}$



Chain of assumptions in ML






deeplearning.ai

Setting up
your goal

Single **number**
evaluation metric

Using a single number evaluation metric



→ Of examples recognized as cost, what % actually are costs?

→ what % of actual costs are correctly recognized

Classifier	Precision	Recall
A	95%	90%
B	98%	85%

F₁ score = "Average" of P and R.

$$\left(\frac{2}{\frac{1}{P} + \frac{1}{R}} \right) \text{ "Harmonic mean"}$$

Dev set + Single number evaluation metric
real speed up iterating

Another example

Algorithm	US	China	India	Other
A	<u>3%</u>	7%	5%	9%
B	5%	6%	5%	10%
C	2%	3%	4%	5%
D	5%	8%	7%	2%
E	4%	5%	2%	4%
F	7%	11%	8%	12%





deeplearning.ai

Setting up
your goal

Satisficing **and**
optimizing metrics

Another cat classification example

Classifier	Accuracy	Running time
A	90%	80ms
B	92%	95ms
C	95%	1,500ms

$$\text{Cost} = \text{accuracy} - 0.5 \times \text{Running Time}$$

maximize accuracy

subject to Running Time \leq 100 ms.

N metrics : 1 optimizing
N-1 satisfying

Wakewords / Trigger words

Alexa, OK Google,

Hey Siri, nihao baidu
你好百度

accuracy.

#false positive

maximize accuracy.

s.t. \leq 1 false positive
every 24 hours.



deeplearning.ai

Setting up
your goal

Train/dev/test
distributions

Cat classification dev/test sets

development set, hold out cross validation set

Regions:

- US
- UK
- Other Europe
- South America
- India
- China
- Other Asia
- Australia


Dev

Test

→ Randomly shuffle into dev/test



dev set
+
metric



True story (details changed)

[Optimizing on dev set on loan approvals for
medium income zip codes

↑

$x \rightarrow y$ (repay loan?)



[Tested on low income zip codes


~ 3 month



Guideline

Choose a dev set and test set to reflect data you expect to get in the future and consider important to do well on.

training



dev
metric

test




deeplearning.ai

Setting up your goal

Size of dev and test sets

Old way of splitting data



Size of dev set

A B

Set your dev set to be big enough to detect differences in
algorithm/models you're trying out.

100 : small
└ 1%

1,000

10,000

100,000


^A 97% → ^B 97.1%
0.1%
└

0.01%
└
0.001%

Online advertising

Size of test set

- Set your test set to be big enough to give high confidence in the overall performance of your system.





deeplearning.ai

Setting up your goal

When to **change**
dev/test sets and
metrics

Cat dataset examples

Metric + Dev : Prefer A
You/users : Prefer B.

→ Metric: classification error

Algorithm A: 3% error

→ pornographic

✓ Algorithm B: 5% error

Error: $\frac{1}{\sum_i w^{(i)}} \cdot \frac{1}{m_{dev}} \sum_{i=1}^{m_{dev}} w^{(i)} \mathbb{I}\{y_{pred}^{(i)} \neq y^{(i)}\}$

↪ $w^{(i)} = \begin{cases} 1 & \text{if } x^{(i)} \text{ is non-porn} \\ 10 & \text{if } x^{(i)} \text{ is porn} \end{cases}$

$\mathbb{I}\{y_{pred}^{(i)} \neq y^{(i)}\}$
predicted value (0/1)

Orthogonalization for cat pictures: anti-porn

- 1. So far we've only discussed how to define a metric to evaluate classifiers. ← Place target ↗
- 2. Worry separately about how to do well on this metric. ↗
- ↖ Aim (shoot at target)

$$\rightarrow J = \frac{1}{\sum w^{(i)}} \sum_{i=1}^m w^{(i)} \mathcal{L}(\hat{y}^{(i)}, y^{(i)})$$



Another example

Algorithm A: 3% error

✓ Algorithm B: 5% error ←

→ Dev/test



→ User images



If doing well on your metric + dev/test set does not correspond to doing well on your application, change your metric and/or dev/test set.




deeplearning.ai

Comparing to human-level performance

Why **human**-level performance?

Comparing to human-level performance



Why compare to human-level performance

Humans are quite good at a lot of tasks. So long as ML is worse than humans, you can:

- - Get labeled data from humans. (x, y)
- - Gain insight from manual error analysis:
Why did a person get this right?
- - Better analysis of bias/variance.




deeplearning.ai

Comparing to human-
level performance


Avoidable bias

Bias and Variance




high bias

underfitting



“just right”



high variance

overfitting


Bias and Variance

Cat classification

Human-level $\approx 0\%$ ----

Training set error:

Dev set error:




high bias

high bias
high variance

low bias
low variance

Cat classification example





deeplearning.ai

Comparing to human-level performance

Understanding
human-level
performance

Human-level error as a proxy for Bayes error

Medical image classification example:

Suppose:

(a) Typical human 3 % error

→ (b) Typical doctor 1 % error

(c) Experienced doctor 0.7 % error

→ (d) Team of experienced doctors .. 0.5 % error ←

Bayes error \leq 0.5 %

What is “human-level” error?



Error analysis example




Human (proxy for Bayes error)

↑ Avoidable bias
↓

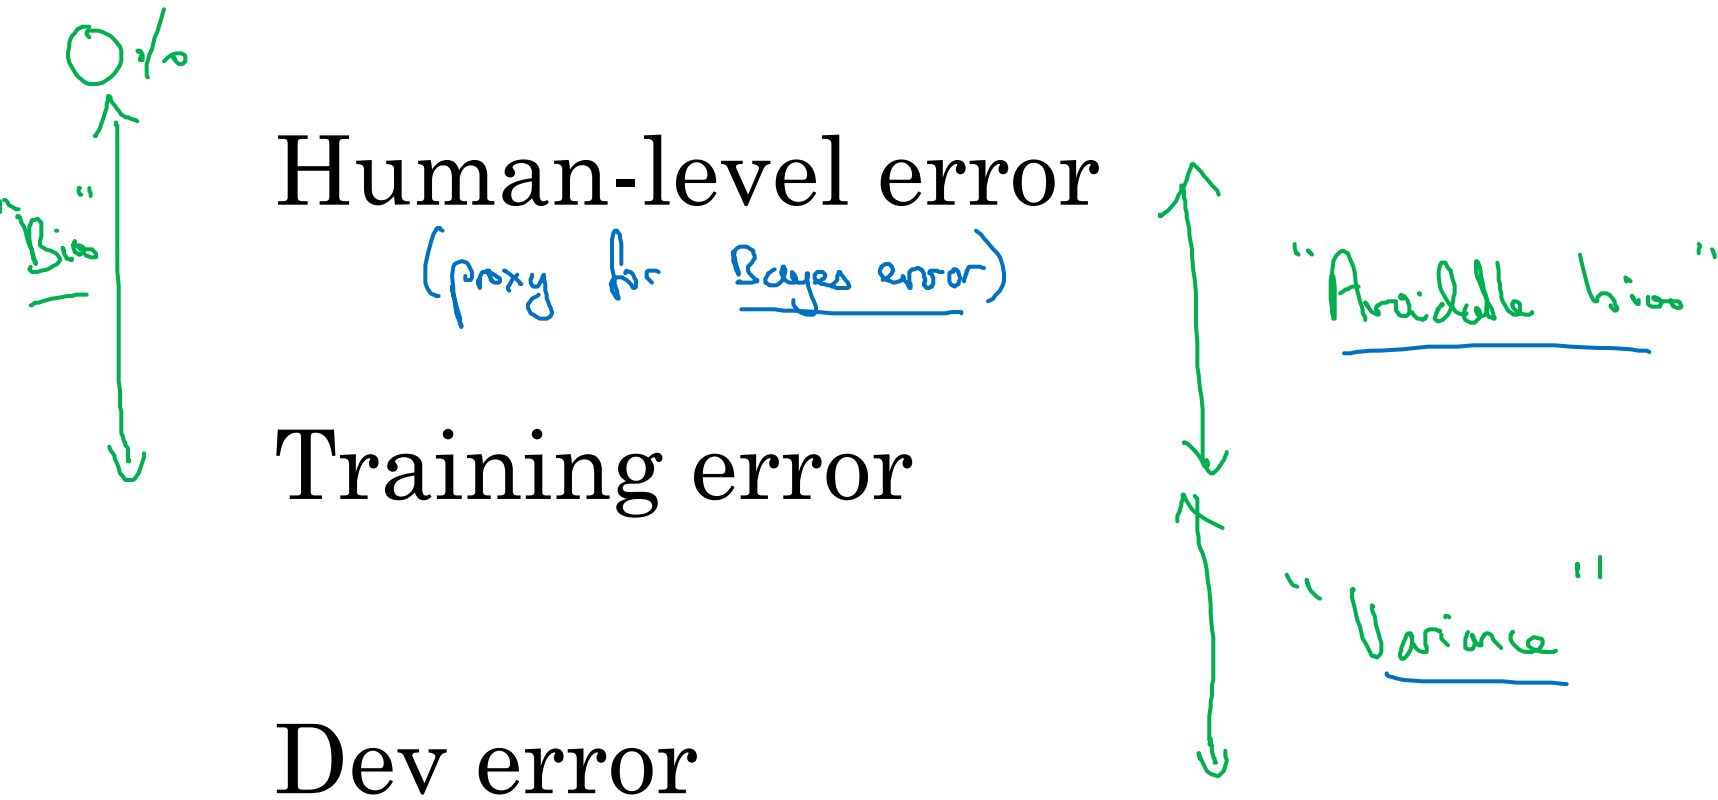
Training error

↑ Variance
↓

Dev error



Summary of bias/variance with human-level performance





deeplearning.ai

Comparing to human-
level performance

Surpassing **human-**
level performance

Surpassing human-level performance

Team of humans

0.5%

One human

0.1 ~~1.0%~~


Training error

0.6%

Dev error

0.2
0.8%

What is avoidable bias?



Problems where ML significantly surpasses human-level performance

- - Online advertising
- - Product recommendations
- - Logistics (predicting transit time)
- - Loan approvals

Structured data

Not natural perception

Lots of data

- Speech recognition
- Some image recognition
- Medical
 - ECG, Skin cancer, ...



deeplearning.ai

Comparing to human-
level performance

Improving your **model**
performance

The two fundamental assumptions of supervised learning

1. You can fit the training set pretty well.



~ Avoidable bias

2. The training set performance generalizes pretty well to the dev/test set.



~ Variance

Reducing (avoidable) bias and variance

Human-level



Avoidable bias

Training error



Variance

Dev error

Train bigger model

Train longer/better optimization algorithms

- momentum, RMSprop, Adam

NN architecture/hyperparameters search

RNN
CNN

More data

Regularization

- L_2 , dropout, data augmentation

NN architecture/hyperparameters search